

Bach Nguyen

bachthebusiness@gmail.com | [linkedin.com/in/ben1203/](https://www.linkedin.com/in/ben1203/) | 417-351-7946 | Willing to relocate

Data Engineer with 2+ years designing ELT pipelines, dimensional data models, and analytics solutions at enterprise scale with strong expertise in PySpark, SQL across cloud platforms (AWS and Azure), processing 50TB+ data across 700+ tables.

SKILLS

Languages: Python, SQL

Data Processing & Visualization: Apache Spark, dbt, Apache Airflow, SSMS, Data Factory, AWS QuickSight, PowerBI, Tableau

Cloud: AWS (S3, Glue, Redshift, Lambda), Microsoft Azure (Fabric, Databricks)

Certification: DP-700 (Microsoft Certified: Fabric Data Engineer Associate - [Credential Link](#))

EXPERIENCE

Analytics Engineer II, CoxHealth – Springfield, MO

Jan. 2026 - Present

- Architected end-to-end Workday-to-Microsoft Fabric ELT platform across a team of 5, processing 50M+ records/day via PySpark, ADF, and SQL; delivered 3 dbt semantic models consumed directly by C-suite and enterprise stakeholders.
- Improved Spark job throughput by 20% across a 2B+ row dataset by diagnosing and resolving data skew through strategic repartitioning, cutting stage-level shuffle overhead.
- Engineered a Workday SOAP API ingestion framework processing \$4B+ in workforce financial data daily, serving as the primary data source for director-level and above reporting across the enterprise.
- Migrated 1.5B+ records from Oracle Cerner and Azure Synapse into Microsoft Fabric, decommissioning 2 legacy systems, enabling 8 enterprise reports, and saving \$15M annually in licensing and infrastructure costs.

Data Engineer, CoxHealth – Springfield, MO

Jan. 2025 - Jan. 2026

- Designed a standardized ELT framework in PySpark, SQL, Airflow, and ADF supporting API and file ingestion with built-in CI/CD, validations, error logging, and metadata mappings, reducing pipeline development time by 75% across the team.
- Decommissioned Snowflake and accelerated project delivery by 30% by building 5 medallion-architecture ELT pipelines ingesting from AWS S3, Oracle, and Fabric, enabling 6 reports via metadata-driven PySpark and Airflow orchestration.
- Reclaimed 60TB+ of storage and saved \$18K annually by building scalable data lake lifecycle management workflows, automating 1TB+ weekly cleanup, and enforcing governance policies across 700+ tables.
- Deployed org-wide AI workflows using Microsoft 365 Copilot and Azure AI Foundry agents (GPT-4.1) integrated into Teams, saving 25+ hours/month across 50 users including directors and C-suite, automating documentation and code review.
- Stabilized legacy pipeline infrastructure by resolving 50+ production tickets through root cause analysis and targeted optimizations, improving reliability and reducing recurring incident rate.

Data Engineer Intern, CoxHealth – Springfield, MO

Aug. 2024 - Dec. 2024

- Deployed an ELT pipeline using PySpark, SQL, and REST APIs migrating 400M+ rows of data from Oracle to Fabric.
- Recovered 2% of lost data in AWS S3 by resolving a migration defect by tracing errors across pipelines and models.

Business Intelligence Intern, CoxHealth – Springfield, MO

Jun. 2024 - Aug. 2024

- Designed a Power BI dashboard integrating 60+ datasets tracking operational performance across 30+ departments.
- Delivered 5+ SQL ad-hoc requests by integrating 800M+ records across Snowflake, Oracle, and Microsoft Fabric.

PROJECTS

Real-Time E-Commerce Event Pipeline - AWS Kinesis, Lambda, S3, Glue, Redshift

- Designed a real-time ELT pipeline ingesting 500K+ simulated daily clickstream events via AWS Kinesis, Lambda, S3 (Bronze), transforming to Parquet with AWS Glue (Silver), and loading a Star Schema into Amazon Redshift (Gold).
- Reduced average query latency by 40% through Redshift distribution key optimization and S3 partition pruning.

EDUCATION

Missouri State University – Bachelor of Science in Data Analytics

Dec. 2024